

[Henrique Yamahata](#)

Titanic Problem, Begginer

last run 5 hours ago · IPython Notebook HTML · 58 views

using data from [Titanic: Machine Learning from Disaster](#) · 👁 Public▲
3

voters

[Notebook](#)[Code](#)[Data \(1\)](#)[Output](#)[Comments \(1\)](#)[Log](#)[Versions \(2\)](#)[Forks \(2\)](#)[Fork Notebook](#)

Notebook

Titanic Problem:

We want to discover which passengers survived through the data.

This notebook is divide by:

- Data analysis
- Feature Engineer at:
 - Gender, Embarked type, Name, Age and Fare
- Modeling with:
 - KNeighborsClassifier, LogisticRegression, DecisionTreeClassifier, RandomForestClassifier
 - Score and Cross-Validation

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

In [2]:

```
titanic = pd.read_csv("../input/train.csv")
titanic_test = pd.read_csv("../input/test.csv")
# Creating a list with two files, more accuracy for to the math to fill NaN values
combined = [titanic, titanic_test]
```

In [3]:

```
titanic.head()
```

Out [3]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

In [4]:

```
titanic_test.head()
```

Out [4]:

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	892	3	Kelly, Mr. James	male	34.5	0	0	330911	7.8292	NaN	Q
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	363272	7.0000	NaN	S
2	894	2	Myles, Mr. Thomas Francis	male	62.0	0	0	240276	9.6875	NaN	Q
3	895	3	Wirz, Mr. Albert	male	27.0	0	0	315154	8.6625	NaN	S
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.0	1	1	3101298	12.2875	NaN	S

Working with Gender

```
In [5]: #Checking if exist some NaN value  
len(titanic[titanic['Sex'].isnull()])
```

```
Out[5]:  
0
```

```
In [6]: #How many unique enters this array have.  
titanic["Sex"].unique()
```

```
Out[6]:  
array(['male', 'female'], dtype=object)
```

```
In [7]: #Checking which gender have more survivors  
titanic[['Survived', 'Sex']].groupby('Sex').mean()
```

```
Out[7]:
```

	Survived
Sex	
female	0.742038
male	0.188908

In [8]:

```
# Replacing Categorical variables by continuous, with this for and this list(combined),
# we can replace Sex in titanic and test_titanic

dicsex = {"male": 0, "female": 1}
for dfsex in combined:
    dfsex['Sex'] = dfsex['Sex'].map(dicsex)

#other method
#titanic.loc[titanic["Sex"] == "male", "Sex"] = 0
#titanic.loc[titanic["Sex"] == "female", "Sex"] = 1
```

In [9]:

```
titanic.head()
```

Out[9]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	0	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	1	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	0	35.0	0	0	373450	8.0500	NaN	S

Working with Embarked type

```
In [10]: #Checking if exist some NaN value
len(titanic[titanic['Embarked'].isnull()])
```

Out[10]:

2

```
In [11]: titanic[titanic['Embarked'].isnull()]
```

Out[11]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
61	62	1	1	Icard, Miss. Amelie	1	38.0	0	0	113572	80.0	B28	NaN
829	830	1	1	Stone, Mrs. George Nelson (Martha Evelyn)	1	62.0	0	0	113572	80.0	B28	NaN

```
In [12]: # Trying to found an insight, connecting all common variables of these people to predict where t
hey embarked, without success
titanic[(titanic['Pclass'] == 1) & (titanic['Survived'] == 1) & (titanic['Sex'] == 1)].groupby(
    'Embarked').sum()
```

Out[12]:

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
Embarked								


C	18038	42	42	42	1320.0	22	13	4943.8208
Q	413	1	1	1	33.0	1	0	90.0000
S	23788	46	46	46	1412.0	27	26	4450.1917

In [13]:

```
# Here we can see, at all classes most of people embarked on "S", so to fill this data with less variation we put "S".  
sns.countplot(x = 'Pclass', data = titanic, hue = 'Embarked')
```

Out[13]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f71d1600cc0>



In [14]:

```
titanic["Embarked"] = titanic["Embarked"].fillna("S")
```

```
In [15]:  
# Same concept to replace used at Sex column  
dic_embarked = {"S": 0, "C": 1, "Q": 2}  
for df_embarked in combined:  
    df_embarked['Embarked'] = df_embarked['Embarked'].map(dic_embarked)
```

Working with Name

Combining both Dataset, BECAUSE EXIST THE POSSIBILITY THAT IN ONE DF DOESN'T EXIST THE SAME PRONOUNS TREATMENT in the other

```
In [16]:  
# getting all Title from Name column in both DataFrames through the list and creating a new column with those titles  
for df in combined:  
    df['Title'] = df['Name'].str.extract(' ([A-Za-z]+)\.', expand=True)
```

```
In [17]:  
# concat both updated dataframes to see the distribution of titles at titanic  
combined_df = pd.concat([titanic, titanic_test], axis = 0)  
combined_df['Title'].value_counts()
```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:2: FutureWarning: Sorting because n on-concatenation axis is not aligned. A future version of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=True'.

To retain the current behavior and silence the warning, pass sort=False

Out[17]:

Mr	757
Miss	260
Mrs	197
Master	61
Dr	8
Rev	8
Col	4
Mlle	2
Major	2
Ms	2
Mme	1
Countess	1
Lady	1
Don	1
Capt	1
Dona	1
Jonkheer	1
Sir	1

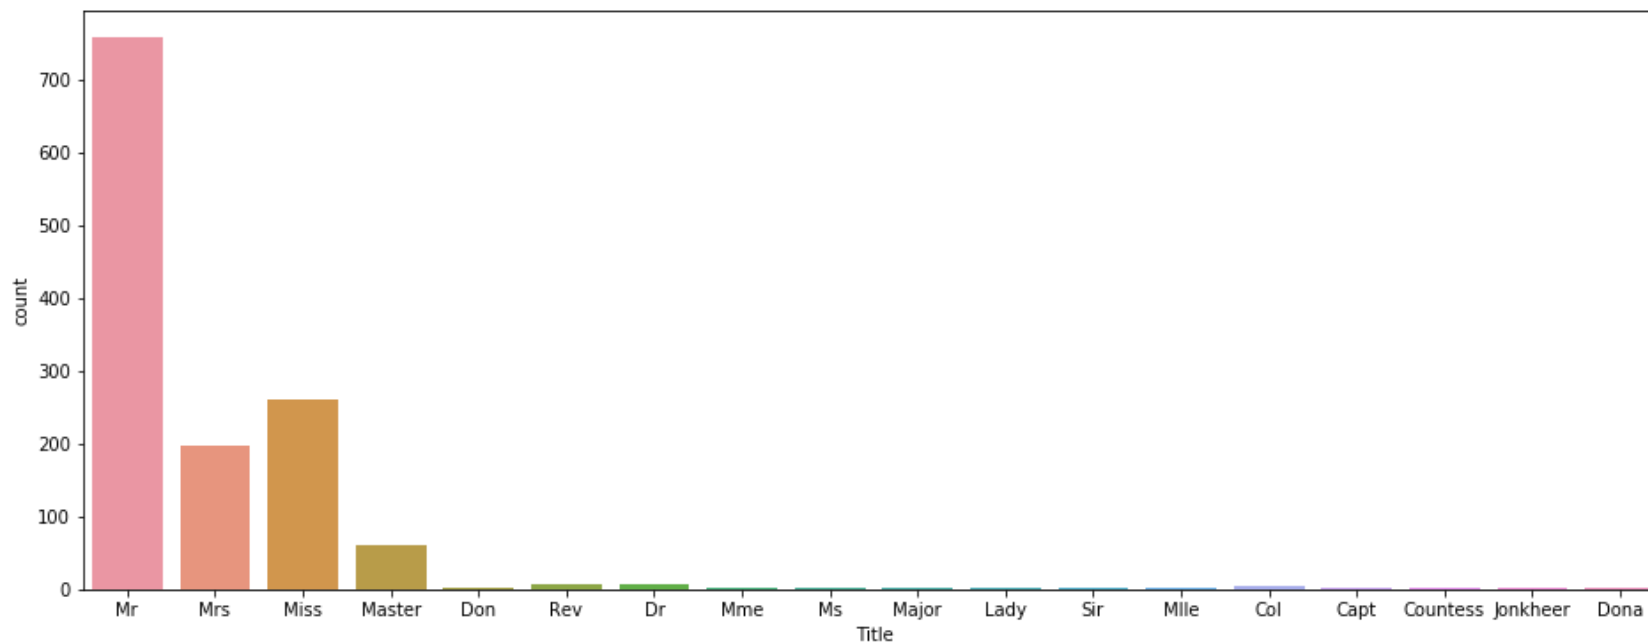
Name: Title, dtype: int64

In [18]:

```
plt.subplots(figsize = (16,6))  
sns.countplot(x = 'Title', data = combined_df)
```

Out[18]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f71d149e940>



In [19]:

```

# Same concept to replace used at Sex column, considering the 4 largest groups of people and the
# rest of them in 1 group
# Mr: 0
# Miss: 1
# Mrs: 2
# Master: 3
# Others: 4
titlemap = {"Mr": 0,
            "Miss": 1,
            "Mrs": 2,
            "Master": 3,
            "Dr": 4, "Rev": 4, "Col": 4, "Major": 4, "Mlle": 4, "Countess": 4, "Ms": 4,
            "Lady": 4, "Jonkheer": 4, "Don": 4, "Dona": 4, "Mme": 4, "Capt": 4, "Sir": 4 }
for df in combined:
    df['Title'] = df['Title'].map(titlemap)

```

In [20]:

```
titanic.head()
```

Out[20]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title
0	1	0	3	Braund, Mr. Owen Harris	0	22.0	1	0	A/5 21171	7.2500	NaN	0	0
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	1	38.0	1	0	PC 17599	71.2833	C85	1	2
2	3	1	3	Heikkinen, Miss. Laina	1	26.0	0	0	STON/O2. 3101282	7.9250	NaN	0	1
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	1	35.0	1	0	113803	53.1000	C123	0	2
4	5	0	3	Allen, Mr. William Henry	0	35.0	0	0	373450	8.0500	NaN	0	0

Working with Age


In [21]:

```
# Lets see how is the distribution by gender for people who survived and whos dont.
survived = titanic[titanic['Survived']==1]['Sex'].value_counts()
# Extract how many peoples for each sex survived
dead = titanic[titanic['Survived']==0]['Sex'].value_counts()
# Extract how many peoples for each sex who not survived
```

```
# Extract how many peoples for each sex not survived
df = pd.DataFrame([survived,dead])
df.columns= ['Male', 'Female']
df.index = ['Survived','Dead']
df.plot(kind='bar',stacked=True, figsize=(10,5))
plt.legend(bbox_to_anchor=(1, 1), loc=2, borderaxespad=1)
```

Out[21]:

<matplotlib.legend.Legend at 0x7f71d1464208>



In [22]:

```
#looking for null values ate Age.
len(titanic[titanic['Age'].isnull()])
```

Out[22]:

177

In [23]:

```
# Here we need to make a decision, which variable we'll relate with Age to fill the empty values
titanic[titanic['Age'].isnull()].head()
```

Out[23]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title
5	6	0	3	Moran, Mr. James	0	NaN	0	0	330877	8.4583	NaN	2	0
17	18	1	2	Williams, Mr. Charles Eugene	0	NaN	0	0	244373	13.0000	NaN	0	0
19	20	1	3	Masselmani, Mrs. Fatima	1	NaN	0	0	2649	7.2250	NaN	1	2
26	27	0	3	Emir, Mr. Farred Chehab	0	NaN	0	0	2631	7.2250	NaN	1	0
28	29	1	3	O'Dwyer, Miss. Ellen "Nellie"	1	NaN	0	0	330959	7.8792	NaN	2	1

In [24]:

```
# looking to the first possibility, calculate new ages through Pclass
combined_df[['Age', 'Pclass']].groupby('Pclass').mean()
```

Out[24]:

	Age
Pclass	
1	39.159930
2	29.506705
3	24.816367

In [25]:

```
#Here we have more accuracy, title it is very related to age  
combined_df = pd.concat([titanic, titanic_test], axis = 0)  
combined_df[['Age', 'Title']].groupby('Title').mean()
```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:2: FutureWarning: Sorting because n on-concatenation axis is not aligned. A future version of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=True'.

To retain the current behavior and silence the warning, pass sort=False

Out[25]:

	Age
Title	
0	32.252151
1	21.774238
2	36.994118
3	5.482642
4	42.656250

In [26]:

```
# this is one function to verify all null values at AGE, and substitute to this respective new value considering Title
# using combined_df, we can be more accurate to get the new mean values, because we consider a bigger data
def impute_age(cols):
    Age = cols[0]
    Title = cols[1]

    if pd.isnull(Age):
        if Title == 0:
            return combined_df['Age'][combined_df['Title'] == 0].mean()
        elif Title == 1:
            return combined_df['Age'][combined_df['Title'] == 1].mean()
        elif Title == 2:
            return combined_df['Age'][combined_df['Title'] == 2].mean()
        elif Title == 3:
            return combined_df['Age'][combined_df['Title'] == 3].mean()
        else:
            return combined_df['Age'][combined_df['Title'] == 4].mean()
    else:
        return Age
```

In [27]:

```
# apply the function on DF's, titanic and titanic_test
titanic['Age'] = titanic[['Age', 'Title']].apply(impute_age, axis=1)
titanic_test['Age'] = titanic_test[['Age', 'Title']].apply(impute_age, axis=1)
```

In [28]:

```
# To improve our machine learning model, we need to smooth our data, so we'll divide our Age values in 5 categories
for df age in combined:
```

```

df_age.loc[ df_age['Age'] <= 16, 'Age'] = 0,
df_age.loc[(df_age['Age'] > 16) & (df_age['Age'] <= 26), 'Age'] = 1,
df_age.loc[(df_age['Age'] > 26) & (df_age['Age'] <= 36), 'Age'] = 2,
df_age.loc[(df_age['Age'] > 36) & (df_age['Age'] <= 62), 'Age'] = 3,
df_age.loc[ df_age['Age'] > 62, 'Age'] = 4
#using this for, we can substitute all values at titanic and also titanic_test

```

In [29]:


```

# Lets see how is the distribution by Age for people who survived and whos dont.
survivedAge = titanic[titanic['Survived']==1]['Age'].value_counts()
# Extract how many peoples for each Age survived
deadAge = titanic[titanic['Survived']==0]['Age'].value_counts()
# Extract how many peoples for each Age not survived
df = pd.DataFrame([survivedAge,deadAge])
df.columns= ['0-16','17-26', '27-36', '37-62', '63+']
df.index = ['Survived','Dead']
df.plot(kind='bar',stacked=True, figsize=(10,5))
plt.legend(bbox_to_anchor=(1, 1), loc=2, borderaxespad=1)

```

Out[29]:

<matplotlib.legend.Legend at 0x7f71d15010f0>






In [30]:

```
df_age2 = titanic[['Age', 'Survived']]
sns.countplot(x = 'Age', hue = 'Survived', data = df_age2)
```

Out[30]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f71d1399e80>



In [31]:

```
titanic_test.head()
```

Out[31]:

PassengerId	Survived	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title
-------------	----------	------	-----	-----	-------	-------	--------	------	-------	----------	-------

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title
0	892	3	Kelly, Mr. James	0	2.0	0	0	330911	7.8292	NaN	2	0
1	893	3	Wilkes, Mrs. James (Ellen Needs)	1	3.0	1	0	363272	7.0000	NaN	0	2
2	894	2	Myles, Mr. Thomas Francis	0	3.0	0	0	240276	9.6875	NaN	2	0
3	895	3	Wirz, Mr. Albert	0	2.0	0	0	315154	8.6625	NaN	0	0
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	1	1.0	1	1	3101298	12.2875	NaN	0	2

Working with FARE

In [32]:

```
#looking for null values ate Fare, Test DF.
titanic_test[titanic_test['Fare'].isnull()]
```

Out[32]:

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title
152	1044	3	Storey, Mr. Thomas	0	3.0	0	0	3701	NaN	NaN	0	0

In [33]:

```
#looking for null values ate Fare, Train DF.
titanic[titanic['Fare'].isnull()]
```

Out[33]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title
--	-------------	----------	--------	------	-----	-----	-------	-------	--------	------	-------	----------	-------

In [34]:

```
# In general, the fare paid is directly relate to the class. the miss value was replaced by the  
mean value at third class  
titanic_test['Fare'] = titanic_test['Fare'].fillna(combined_df['Fare'][combined_df['Pclass'] ==  
3].mean())
```

In [35]:

```
combined_df = pd.concat([titanic, titanic_test], axis = 0)
```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:1: FutureWarning: Sorting because n
on-concatenation axis is not aligned. A future version
of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=True'.

To retain the current behavior and silence the warning, pass sort=False

"""Entry point for launching an IPython kernel.

In [36]:

```
combined_df['Fare'].describe()  
# Isn't good to see the distribution in that way
```

Out[36]:

count	1309.000000
mean	33.280206
std	51.741830
min	0.000000

```

25%          7.895800
50%          14.454200
75%          31.275000
max          512.329200
Name: Fare, dtype: float64

```

In [37]:

```

# To improve our machine learning model, we need to smooth our data, so we'll divide our Fare va
lues in 4 categories
for dataset in combined:
    dataset.loc[ dataset['Fare'] <= 17, 'Fare'] = 0,
    dataset.loc[(dataset['Fare'] > 17) & (dataset['Fare'] <= 30), 'Fare'] = 1,
    dataset.loc[(dataset['Fare'] > 30) & (dataset['Fare'] <= 100), 'Fare'] = 2,
    dataset.loc[ dataset['Fare'] > 100, 'Fare'] = 3

```

In [38]:

```
titanic_test.head()
```

Out[38]:

	PassengerId	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title
0	892	3	Kelly, Mr. James	0	2.0	0	0	330911	0.0	NaN	2	0
1	893	3	Wilkes, Mrs. James (Ellen Needs)	1	3.0	1	0	363272	0.0	NaN	0	2
2	894	2	Myles, Mr. Thomas Francis	0	3.0	0	0	240276	0.0	NaN	2	0
3	895	3	Wirz, Mr. Albert	0	2.0	0	0	315154	0.0	NaN	0	0
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	1	1.0	1	1	3101298	0.0	NaN	0	2

Cabin

In [39]:

```
for dfcabin in combined:
    dfcabin['Cabin'] = dfcabin['Cabin'].str[:1]
```

In [40]:

```
# if we try to associate Cabin location to Fare, data are very confused
print(titanic[titanic['Fare'] == 0]['Cabin'].unique())
print(titanic[titanic['Fare'] == 1]['Cabin'].unique())
print(titanic[titanic['Fare'] == 2]['Cabin'].unique())
print(titanic[titanic['Fare'] == 3]['Cabin'].unique())
```

```
[nan 'G' 'D' 'F' 'E' 'B' 'A']
[nan 'C' 'F' 'D' 'B' 'A' 'E']
['C' 'E' nan 'A' 'D' 'B' 'F' 'T']
['C' 'B' 'D' nan 'E']
```


In [41]:

```
# if we try to associate Cabin location to Pclass, we can see a pattern
print(titanic[titanic['Pclass'] == 1]['Cabin'].unique())
print(titanic[titanic['Pclass'] == 2]['Cabin'].unique())
print(titanic[titanic['Pclass'] == 3]['Cabin'].unique())
```

```
['C' 'E' 'A' nan 'B' 'D' 'T']
[nan 'D' 'F' 'E']
[nan 'G' 'F' 'E']
```

In [42]:

```
PC1 = titanic[titanic['Pclass'] == 1]['Cabin'].value_counts()
PC2 = titanic[titanic['Pclass'] == 2]['Cabin'].value_counts()
PC3 = titanic[titanic['Pclass'] == 3]['Cabin'].value_counts()
dfCabin = pd.DataFrame([PC1, PC2, PC3])
dfCabin.index = ['1 Class', '2 Class', '3 Class']
dfCabin.plot(kind = 'bar', stacked = True, figsize = (12,6))
plt.style.use('bmh')
```



In [43]:

```
# Due Cabin A,B,T,C only exist at First class, they become only A
titanic['Cabin'].replace(['B', 'T', 'C'], ['A', 'A', 'A'], inplace = True);
titanic_test['Cabin'].replace(['B', 'T', 'C'], ['A', 'A', 'A'], inplace = True);
```

```
In [44]: titanic['Cabin'].unique()
```

```
Out[44]: array([nan, 'A', 'E', 'G', 'D', 'F'], dtype=object)
```

```
In [45]: dicCabins = {"A": 0, "D": 0.5, "E": 1, "F": 1.5, "G": 2}
for dataset2 in combined:
    dataset2['Cabin'] = dataset2['Cabin'].map(dicCabins)
```

```
In [46]: def impute_cabin(cols):
        Cabin = cols[0]
        Pclass = cols[1]

        if pd.isnull(Cabin):
            if Pclass == 1:
                return 0
            elif Pclass == 2:
                return 1
            else:
                return 1.5
        else:
            return Cabin

titanic['Cabin'] = titanic[['Cabin', 'Pclass']].apply(impute_cabin, axis=1)
titanic_test['Cabin'] = titanic[['Cabin', 'Pclass']].apply(impute_cabin, axis=1)
```

MODELING

In [47]:

```
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
k_fold = KFold(n_splits=100, shuffle=True, random_state=0)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_teste, Y_train, Y_teste = train_test_split(titanic.drop(["Survived", 'PassengerId', 'Name', 'Ticket', 'Cabin'], axis=1),
titanic["Survived"], test_size = 0.3, random_state = 101)
```

In [48]:

```
X_train = titanic.drop(["Survived", 'PassengerId', 'Name', 'Ticket'], axis=1)
Y_train = titanic["Survived"]
X_test  = titanic_test.drop(["PassengerId", 'Name', 'Ticket'], axis=1)
```

In [49]:

```
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors = 25)
knn.fit(X_train, Y_train)
Y_pred = knn.predict(X_test)
acc_knn = round(knn.score(X_train, Y_train) * 100, 2)
score_knn = cross_val_score(knn, X_train, Y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy')
print('KNN Cross: {} \n KNN: {}'.format(round(np.mean(score_knn)*100,3), acc_knn))
```

```
KNN Cross: 82.167
```

```
KNN:      83.05
```


In [50]:

```
from sklearn.linear_model import LogisticRegression
logistic = LogisticRegression()
logistic.fit(X_train, Y_train)
Y_pred = logistic.predict(X_test)
acc_log = round(logistic.score(X_train, Y_train) * 100, 2)
score_lr = cross_val_score(logistic, X_train, Y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy'
)
print('Logistic Cross: {} \n Logistic:      {}'.format(round(np.mean(score_lr)*100,2), acc_log))
```

Logistic Cross: 81.17

Logistic: 81.71

In [51]:

```
from sklearn.tree import DecisionTreeClassifier

decision_tree = DecisionTreeClassifier()
decision_tree.fit(X_train, Y_train)
Y_pred = decision_tree.predict(X_test)
acc_decision_tree = round(decision_tree.score(X_train, Y_train) * 100, 2)
score_dt = cross_val_score(decision_tree, X_train, Y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy')
print('Decision Tree Cross: {} \n Decision Tree:      {}'.format(round(np.mean(score_dt)*100,2),
acc_decision_tree))
```

Decision Tree Cross: 79.72

Decision Tree: 89.79

In [52]:

```
from sklearn.ensemble import RandomForestClassifier
```

```
random_forest = RandomForestClassifier(n_estimators=200)
random_forest.fit(X_train, Y_train)
Y_pred = random_forest.predict(X_test)
random_forest.score(X_train, Y_train)
acc_random_forest = round(random_forest.score(X_train, Y_train) * 100, 2)
score_rf = cross_val_score(random_forest, X_train, Y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy')
print('Random Forest Cross: {} \n Random Forest:      {}'.format(round(np.mean(score_rf)*100,2),
acc_random_forest))
```

Random Forest Cross: 80.26

Random Forest: 89.79

In [53]:

```
from sklearn.ensemble import GradientBoostingClassifier

gbk = GradientBoostingClassifier()
gbk.fit(X_train, Y_train)
Y_pred = gbk.predict(X_test)
acc_gbk = round(gbk.score(X_train, Y_train) * 100, 2)
score_gbk = cross_val_score(gbk, X_train, Y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy')
print('Gradient Boosting Classifier Cross: {} \n Gradient Boosting Classifier:      {}'.format(round(np.mean(score_gbk)*100,2), acc_gbk))
```

Gradient Boosting Classifier Cross: 82.46

Gradient Boosting Classifier: 86.08

In [54]:

```
from sklearn.svm import SVC
```

```

svc = SVC(gamma = 'scale')
svc.fit(X_train, Y_train)
Y_pred = svc.predict(X_test)
acc_svc = round(svc.score(X_train, Y_train) * 100, 2)
score_svc = cross_val_score(svc, X_train, Y_train, cv=k_fold, n_jobs=1, scoring = 'accuracy')
print('SVC Cross: {} \n SVC:      {}'.format(round(np.mean(score_svc)*100,2), acc_svc))

```

SVC Cross: 82.28

SVC: 83.95

In [55]:


```

from sklearn.model_selection import train_test_split
x_trainG, x_testeG, y_trainG, y_testeG = train_test_split(titanic.drop(["Survived", 'PassengerId'
,
                                                                    'Name', 'Ticket', 'Cabin'],
                                                                    axis=1), titanic["Survived"
],
                                                                    test_size = 0.4, random_sta
te = 101)
error_rate = []
for i in range(1,35):
    knnG = KNeighborsClassifier(n_neighbors = i)
    knnG.fit(x_trainG, y_trainG)
    y_predG = knnG.predict(x_testeG)
    error_rate.append(np.mean(y_predG!=y_testeG))
plt.figure(figsize = (14, 8))
plt.plot(range(1, 35), error_rate, color = 'blue', ls = 'dashed', marker = 'o')

```

Out[55]:

[<matplotlib.lines.Line2D at 0x7f71d121cf98>]



In [56]:

```
models = pd.DataFrame({
    'Model': ['Logistic Regression', 'Decision Tree',
              'Random Forest', 'KNN', 'Gradient Boosting Classifier', 'SVC'],
    'Score': [np.mean(score_lr)*100, np.mean(score_dt)*100, np.mean(score_rf)*100,
              np.mean(score_knn)*100, np.mean(score_gbk)*100, np.mean(score_svc)*100])
print('Cross Validation')
models.sort_values(by='Score', ascending=False)
```

Cross Validation

Out[56]:

	Model	Score
4	Gradient Boosting Classifier	82.458333
5	SVC	82.277778
3	KNN	82.166667
0	Logistic Regression	81.166667
2	Random Forest	80.263889
1	Decision Tree	79.722222

In [57]:

```
models = pd.DataFrame({
    'Model': ['Logistic Regression', 'Decision Tree',
              'Random Forest', 'KNN', 'Gradient Boosting Classifier', 'SVC'],
    'Score': [acc_log, acc_decision_tree, acc_random_forest,
              acc_knn, acc_gbk, acc_svc]})
print('Score')
models.sort_values(by='Score', ascending=False)
```

Score

Out[57]:

	Model	Score
1	Decision Tree	89.79
2	Random Forest	89.79
4	Gradient Boosting Classifier	86.08
5	SVC	83.95

3	KNN	83.05
0	Logistic Regression	81.71

In [58]:

```
# Run the Model First
submission = pd.DataFrame({
    "PassengerId": titanic_test["PassengerId"],
    "Survived": Y_pred
})
submission.to_csv('submission.csv', index=False)
```

Did you find this Kernel useful?

Show your appreciation with an upvote


3

k



Comments (1)

All Comments ▼

Sort by

Hotness ▼



Click here to enter a comment...

**Henrique Yamahata**

Kernel Author

• Posted on Latest Version • 37 minutes ago • Options • Reply



0



i'll work with the Tickets and family and improve the code soon !!